

Programming Assignment 3: Membership Service

Due Monday, Oct. 21, 9:00 pm

1 Lab Task

In this lab, you will implement a membership service similar with the one described in slides 17–21 from Week 4. We will refer to all servers as peers.

Each peer starts by knowing a list representing the exhaustive set of hosts that can be part of the group. You can assume that the leader is the first host in the file and that the peers will know who the leader is from the file at startup. Peer ids will be assigned based on the line in the file, starting at 1.

The goal is for peers to build and maintain a list of all alive peers. A membership consists of a unique id, referred to as *view id* and the list of the alive peers. The view id should be monotonically increasing and create a total order, even when the leader changes. Each peer should maintain a view id and a membership list.

Every time the membership list changes each peer should print the new view id and the new list of peers, the list should be printed in increasing order of peer id.

You can assume that the maximum number of peers is 10. You should test your code with minimum 5 nodes and provide the 4 testcases described below.

You can use C/C++, Java, or Go.

1.1 PART 1: Add a peer to the membership list (30 points)

When a peer starts it will contact the leader by sending a JOIN message. The leader will start a protocol with existing members as follows:

- Leader sends a REQ message containing a request id, current view id, operation type, and the id of the peer to be added, to all the other peers in the existing membership list. Operation type is ADD.
- Each peer saves the operation he must perform (request id, current view id, operation type, peer id) and sends back an OK message containing the request id and the current view id.
- When receiving OK messages from all the alive peers (with matching request id and current view id), the leader will increment its view id, add the new peer to its membership list and send a NEWVIEW message that contains the new view id and new membership list to all the members including the new peer.
- When receiving the NEWVIEW message all processes update their view id and membership list and print the new view (view id and membership list).

You can assume that no peer joins before the previous one finished and that there are no crashes. Leader starts first with the membership list containing just himself.

You can optimize the protocol described above, for example send the entire list only to the new peer, and to the other peers just the view id.

Communication can be implemented with either TCP or UDP. Communication should be reliable, if you are using UDP you need to implement a reliable channel.

TESTCASE 1: Each process joins one by one, till all the processes from the configuration file joined. Leader starts first. At the end you should see for each peer how the membership list and view id changed, every time a new peer joined.

OUTPUT: The following output should be printed *verbatim*. Text enclosed in angle brackets (<>) should be replaced with actual values, e.g., [**<COLLEGE>**] → [**<khoury>**]. Every time a peer joins, all peers in the membership should print:

```
{peer_id:<ID>, view_id: <VIEW_ID>, leader: <LEADER_ID>,  
memb_list: [<COMMA_SEPARATED_MEMBERS>]}
```

Where **peer_id** is the ID of the local peer, **view_id** is the current view ID of the local peer, and **leader** is the current leader ID. Example:

```
{peer_id: 2, view_id: 3, leader: 1, memb_list: [1,2,3,5]}
```

1.2 PART 2: Failure detector (10 points)

Each peer will implement a failure detector as follows.

- Each peer will broadcast a HEARTBEAT message every time a timeout T expires. (You can implement the broadcast as a true broadcast or multiple unicasts.)
- When a peer P does not hear from another peer Q within a timeframe that corresponds to twice the period with which HEARTBEAT messages are sent, it will declare Q dead and should print the failure message described below. (Failure of a peer here means missing two consecutive heartbeats from that peer).

All the HEARTBEAT messages should be implemented with UDP. Heartbeat messages are not reliable.

TESTCASE 2: Crash a peer and have all other peers detect its failure. All peers should print on the screen the failure detection message as described below.

OUTPUT: The following output should be printed *verbatim*. When a peer crashes, that peer should print the following before going down:"

```
{peer_id:<ID>, view_id: <VIEW_ID>, leader: <LEADER_ID>, message:"crashing"}
```

All the peers still in the membership should print

```
{peer_id:<ID>, view_id: <VIEW_ID>, leader: <LEADER_ID>,  
message:"peer <PEER_ID> unreachable"}
```

If the peer that crashed was the leader, the following should be printed:

```
{peer_id:<ID>, view_id: <VIEW_ID>, leader: <LEADER_ID>,  
message:"peer <PEER_ID> (leader) unreachable"}
```

1.3 PART 3: Delete a peer from the membership list (20 points)

Once the leader has detected that a process crashed, it will initiate a protocol to remove the member from the list – following a similar process with the one for adding a new member.

- Leader send a REQ message containing request id, current view id, operation type, and peer id of the peer to be removed. Operation type is DEL.
- All peers save the operation they must perform (request id, current view id, operation type, peer id) and send OK with matching request id and current view id.
- Once the leader receives matching OKs from all alive peers (request id and current view id match with the request), the leader will increment the view id and make the change to the list, and send to all peers a NEWVIEW message that contains the new view id and the corresponding updated list. All peers update their view id and membership list and print the new view id and membership list on the screen.

As in the case of join, you can assume that only one peer leaves at the time and that finishing the protocol to update the list will not be interrupted by another event. The leader does not leave or crash. Only the leader needs to detect the crashed peer to update the list, the other processes even if they detect the failures, still need to wait for the leader to initiate the change.

Communication should be implemented as in Part 1.

TESTCASE 3: After all the peers are part of the membership, start crashing them one by one, wait for each view change to finish before crashing another process, In the end you should be left only with the leader.

You can optimize the delete if you want, but the peers will make the changes permanent to the list, i.e. modify the view id and the list only when receiving the NEWVIEW message from the leader.

OUTPUT: When a peer is deleted, the following should be printed (it's the same as for adding).

```
{peer_id:<ID>, view_id: <VIEW_ID>, leader: <LEADER_ID>,  
memb_list: [<COMMA_SEPARATED_MEMBERS>]}
```

1.4 PART 4: Leader failure (30 points)

Once the leader crashed, the peer with the lowest id becomes the new leader, all members should automatically know from the membership list who is the new leader. The new leader initiates a reconciliation phase as follows:

- The new leader sends a NEWLEADER message to all peers asking if they have any pending operations – operations that they have saved when they received a request for changes from the previous leader, but they did not apply to the list yet because the leader died before sending the NEWVIEW message. NEWLEADER message should contain request id, current view id, operation type; operation type is PENDING.
- The peers will respond with a message that contains that request, current view id, the operation, ADD or DEL, and the id of the peer to be added or deleted – if they have any pending. If they do not have anything pending, the operation will be NOTHING.

- The new leader will finish the received operation (remember that we assume that there is only one event being handled and no cascading events) by restarting the add or delete protocol described in Part 1 or Part 3.

You can assume that the change was only one event, either an add or a delete. Communication is as in Part 1 and Part 3. You can optimize the protocol if you want.

TESTCASE 4: crash the current leader once he sent a request for removing a member, but before all processes received that request, specifically the peer that will be the new leader has not received this information. (Remember that the new leader will be the peer with the lowest id). Your testcase should show that the new leader finishes the operation correctly and the pending peer is removed from the list and all peers correctly update the view id and the membership list.

OUTPUT: Messages corresponding to the previous test cases should be printed.

IMPORTANT: You do not need to handle more than one event, one process joins, or one leaves. You do not need to handle "cascading events" – i.e. all joins and leaves finish before another event starts. The only exception is the testcase for the leader failure as described above.

The REQ, OK, NEWVIEW, NEWLEADER, JOIN are message types. ADD, DEL, PENDING are operation types. Do not use strings for them.

2 Implementation

You need to implement this algorithm in C/C++, Java, or Go, and your implementation must allow the user to configure the execution of the process. You will again be using Docker and Docker Compose to package your program. You will be provided with a `docker-compose.yml` file for each testcase. Your program should be able to use that file without modification.

IMPORTANT: all messages must be printed to `stderr` and all printed messages must be terminated with a newline (`\n`). This is needed for our grading scripts.

Note: you must implement all CLI arguments *exactly* as written below to receive full credit (they are important for our grading scripts).

Building:

```
docker build . -t prj3
```

Running this in your project's directory should build your project's Docker image using your Dockerfile. It should copy over the code and hostsfile to the image and compile your project.

Usage:

```
docker run --name <hostname> --network <network> --hostname <hostname> \
    prj3 -h <hostfile> [-d <delay>] [-c delay] [-t]
```

`--name <hostname>`

This specifies the name of the Docker container, this name should match the one in your hostfile

`--network <network>`

This is the user-defined Docker network your project

will run in, refer to tutorial in additional instructions to learn how to create one and why you need one.

`--hostname <hostname>`

This specifies the hostname of the Docker container, this name should match the one in your hostfile. It can be used by other containers in the same network.

`-h hostsfile`

The hostsfile is the path to a file that contains the list of hostnames that the processes are running on. It assumes that each host is running only one instance of the process.

...

All the processes will listen on the same port.

The line number indicates the identifier of the process which starts at 1.

`-d delay`

This delay is a sleep that should occur at the start of the program. It should sleep for “delay” seconds before starting any aspects of the protocol.

`-c delay`

This delay is a sleep that should start immediately after sending an JOIN message. When the sleep ends, the peer should crash (exit).

`-t`

This message is for executing TESTCASE 4. It should be passed only to the starting leader. It causes the leader to execute the following sequence: 1) wait for all hosts in the hostsfile to join (i.e., a NEWVIEW has been sent after each of them joined, in turn), 2) send a REQ message (of type DEL) to all hosts except the next leader, 3) crash (exit).

Orchestration:

```
docker compose -f [PATH_TO_COMPOSE_FILE] up
```

3 Submission Instructions

Your submission must include the following files (10 points):

1. The **SOURCE** and **HEADER** files (no object files or binary)
2. A **MAKEFILE** to compile and to clean your project
3. A **README** file containing your name, instructions to run your code and anything you would like us to know about your program (like errors, special conditions, etc.)
4. A **REPORT** describing the system architectures, state diagrams, design decisions, and implementation issues

Your submission should also include scripts, code, explanations on how to run the 4 TESTCASES described above. To receive full credit for each part, you need to provide the required

testcases.

Submission is through gradescope.

4 Additional resources

You may find the following resources helpful

- Socket programming: <http://beej.us/guide/bgnet/>
- Unix programming links: <http://www.cse.buffalo.edu/~milun/unix.programming.html>
- C/C++ programming link: <http://www.cplusplus.com/>
- Docker tutorial: <https://github.com/iowaguy/docker-tutorial>